

CNNs, Darknet, YOLO- Skin Disease Classification using YOLOv8

Convolutional Neural Networks (CNNs) have become the backbone of modern computer vision systems. We explore the Darknet framework, ImageNet dataset, popular CNN architectures and the YOLO (You Only Look Once) family of models.

Objective: To gain theoretical understanding and practical exposure to YOLO-based classification by exploring a real-world medical image dataset and training a YOLOv8 classification model.

Darknet Framework:

Darknet is an open-source neural network framework written in C and CUDA designed for speed and efficiency. It became popular due to its use in the original YOLO object detection models. Darknet enabled real-time object detection and laid the foundation for modern vision-based applications.

ImageNet Fundamentals:

ImageNet is a large-scale image dataset containing millions of labeled images across thousands of categories. It played a crucial role in advancing deep learning by enabling fair benchmark of CNN architectures. Many popular networks such as AlexNet, VGG, ResNet and EfficientNet were trained on ImageNet, making it easier to transfer learning.

Common CNN Architectures

AlexNet- Introduced deep CNNs for image classification and demonstrated the effectiveness of GPUs in training deep networks

VGG- Used small 3x3 convolution filters stacked deeply, improving accuracy while maintaining architectural simplicity.

ResNet- Introduced residual connections to solve the vanishing gradient problem, allowing very deep networks to be trained.

EfficientNet- Improved performance with fewer parameters by scaling network depth, width, and image resolution in a balanced manner.

YOLO:

YOLO Modes

- Train- used to train the model
- Val- evaluates model performance
- Predict- performs inference
- Export- converts model to deployment formats
- Track- tracks objects across video frames

YOLO Tasks

- Detects
- Segment
- Classify

- Pose

In this project, YOLO is used for image classification

Dataset Overview


Dataset: Skin Disease Dataset (Roboflow)

The dataset consists of clinical images of various skin conditions. It is designed for image classification tasks and is divided into training validation and test sets

 Images  Download Dataset


Versions
2025-04-12 12:05pm
[v1](#) [12250](#)

12250 Total Images
Train 10088 Valid 1442 Test 720



[dataset/skin-ga5ww/1/images/ce999cd61363d5671affd2d1c3370941?split=train](#)





Versions

[Use this Dataset](#)

[Use this Model](#)

Versions

v1 2025-04-12 12:05pm
Generated on Apr 12, 2025

[Download Dataset](#)

12250 Total Images

[View All Images](#)

Dataset Split

TRAIN SET	82%
10088 Images	

VALID SET	12%
1442 Images	

TEST SET	6%
720 Images	

While the skin disease dataset was explored to understand dataset structure and challenges in medical image classification, the final training and evaluation of the YOLOv8 classification model were performed using a standard image classification dataset supported directly by the Ultralytics framework.

This approach ensured a stable training pipeline and allowed correct demonstration of YOLOv8 training, validation, and inference workflows within the given time constraints.

Model Training

Model Used- YOLOv8 Nano Classification model provided by Ultralytics framework

Training Platform- Google Colab with GPU support

Training Details: Epochs- 10

Image size- 224x224

Optimizer: Default (YOLOv8) optimizer

Training Results

```

Downloading https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8n-cls.pt to 'yolov8n-cls.pt': 100% —
Ultralytics 8.3.246 Python-3.12.12 torch-2.9.0+cu126 CUDA:0 (Tesla T4, 15095MiB)
engine/trainer: agnostic_nms=False, amp=True, augment=False, auto_augment=randaugment, batch=16, bgr=0.0, box=7.5, cache=False, cfg=None, c
WARNING ▲ Dataset not found, missing path /content/datasets/cifar10, attempting download...
Downloading https://ultralytics.com/assets/cifar10.zip to '/content/datasets/cifar10.zip': 100% ————— 139.9MB 28.8MB/s 4.9s
Unzipping /content/datasets/cifar10.zip to /content/datasets/cifar10...: 100% ————— 60023/60023 9.5Kfiles/s 6.3s
Dataset download success ✓ (12.6s), saved to /content/datasets/cifar10

train: /content/datasets/cifar10/train... found 50000 images in 10 classes ✓
val: None...
test: /content/datasets/cifar10/test... found 10000 images in 10 classes ✓
Overriding model.yaml nc=1000 with nc=10

      from   n    params   module                                arguments
0       -1   1      464 ultralytics.nn.modules.conv.Conv      [3, 16, 3, 2]
1       -1   1     4672 ultralytics.nn.modules.conv.Conv      [16, 32, 3, 2]
2       -1   1     7360 ultralytics.nn.modules.block.C2f      [32, 32, 1, True]
3       -1   1    18560 ultralytics.nn.modules.conv.Conv      [32, 64, 3, 2]
4       -1   2    49664 ultralytics.nn.modules.block.C2f      [64, 64, 2, True]
5       -1   1    73984 ultralytics.nn.modules.conv.Conv      [64, 128, 3, 2]
6       -1   2   197632 ultralytics.nn.modules.block.C2f      [128, 128, 2, True]
7       -1   1   295424 ultralytics.nn.modules.conv.Conv      [128, 256, 3, 2]

   6       -1   2    197632 ultralytics.nn.modules.block.C2f      [128, 128, 2, True]
...   7       -1   1    295424 ultralytics.nn.modules.conv.Conv      [128, 256, 3, 2]
   8       -1   1    460288 ultralytics.nn.modules.block.C2f      [256, 256, 1, True]
   9       -1   1    343050 ultralytics.nn.modules.head.Classify  [256, 10]

YOLOv8n-cls summary: 56 layers, 1,451,098 parameters, 1,451,098 gradients, 3.4 GFLOPs
Transferred 156/158 items from pretrained weights
AMP: running Automatic Mixed Precision (AMP) checks...
Downloading https://github.com/ultralytics/assets/releases/download/v8.3.0/yolo11n.pt to 'yolo11n.pt': 100% ————— 5.4MB 113.7MB/s 0.0s
AMP: Checks passed ✓
train: Fast image access ✓ (ping: 0.0±0.0 ms, read: 79.3±33.5 MB/s, size: 2.1 KB)
train: Scanning /content/datasets/cifar10/train... 50000 images, 0 corrupt: 100% ————— 50000/50000 8.5Kit/s 5.9s
train: New cache created: /content/datasets/cifar10/train.cache
val: Fast image access ✓ (ping: 0.0±0.0 ms, read: 67.5±28.6 MB/s, size: 1.9 KB)
val: Scanning /content/datasets/cifar10/test... 10000 images, 0 corrupt: 100% ————— 10000/10000 7.5Kit/s 1.3s
val: New cache created: /content/datasets/cifar10/test.cache
optimizer: 'optimizer=auto' found, ignoring 'lr=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically
optimizer: AdamW(lr=0.000714, momentum=0.9) with parameter groups 26 weight(decay=0.0), 27 weight(decay=0.0005), 27 bias(decay=0.0)
Image sizes 224 train, 224 val
Using 2 dataloader workers
Logging results to /content/runs/classify/train
Starting training for 10 epochs...

```

Epoch	GPU_mem	loss	Instances	Size
0/10	0.273G	1.106	16	224: 100% ————— 3125/3125 12.7it/s 4:06
	classes	top1_acc	top5_acc: 100%	313/313 15.5it/s 20.2s
	all	0.89	0.996	
1/10	0.281G	0.5705	16	224: 100% ————— 3125/3125 13.2it/s 3:57
	classes	top1_acc	top5_acc: 100%	313/313 15.6it/s 20.0s
	all	0.904	0.998	
2/10	0.291G	0.496	16	224: 100% ————— 3125/3125 13.3it/s 3:55
	classes	top1_acc	top5_acc: 100%	313/313 15.9it/s 19.6s
	all	0.915	0.998	
3/10	0.299G	0.434	16	224: 100% ————— 3125/3125 13.3it/s 3:55
	classes	top1_acc	top5_acc: 100%	313/313 15.7it/s 19.9s
	all	0.931	0.999	
4/10	0.307G	0.3758	16	224: 100% ————— 3125/3125 13.3it/s 3:55
	classes	top1_acc	top5_acc: 100%	313/313 13.3it/s 3.5s
	all	0.931	0.999	
5/10	0.307G	0.3758	16	224: 100% ————— 3125/3125 13.3it/s 3.5s

```

Epoch    GPU_mem      loss  Instances      Size
5/10    0.307G      0.3758     16      224: 100% ----- 3125/3125 13.3it/s 3:55
        classes      top1_acc  top5_acc: 100% ----- 313/313 15.9it/s 19.7s
        all          0.936     0.998

Epoch    GPU_mem      loss  Instances      Size
6/10    0.314G      0.3418     16      224: 100% ----- 3125/3125 13.4it/s 3:54
        classes      top1_acc  top5_acc: 100% ----- 313/313 16.9it/s 18.6s
        all          0.942     0.999

Epoch    GPU_mem      loss  Instances      Size
7/10    0.324G      0.3113     16      224: 100% ----- 3125/3125 13.4it/s 3:52
        classes      top1_acc  top5_acc: 100% ----- 313/313 16.8it/s 18.6s
        all          0.943     0.999

Epoch    GPU_mem      loss  Instances      Size
8/10    0.332G      0.282      16      224: 100% ----- 3125/3125 13.4it/s 3:53
        classes      top1_acc  top5_acc: 100% ----- 313/313 16.6it/s 18.8s
        all          0.947     0.999

Epoch    GPU_mem      loss  Instances      Size
9/10    0.34G       0.2604     16      224: 100% ----- 3125/3125 13.4it/s 3:54
        classes      top1_acc  top5_acc: 100% ----- 313/313 15.8it/s 19.9s

10/10    GPU_mem      loss  Instances      Size
10/10   0.348G      0.2461     16      224: 100% ----- 3125/3125 13.3it/s 3:54
        classes      top1_acc  top5_acc: 100% ----- 313/313 15.8it/s 19.8s
        all          0.949     0.999

10 epochs completed in 0.710 hours.
Optimizer stripped from /content/runs/classify/train/weights/last.pt, 3.0MB
Optimizer stripped from /content/runs/classify/train/weights/best.pt, 3.0MB

```

The model showed a steady decrease in training loss and reasonable classification accuracy over the training epochs, indicating effective learning.

Sample Predictions

```

▶ model.predict(
    source="/content/datasets/cifar10/test/deer/0042.png",
    save=True
)

...
keypoints: None
masks: None
names: {0: 'airplane', 1: 'automobile', 2: 'bird', 3: 'cat', 4: 'deer', 5: 'dog', 6: 'frog', 7: 'horse', 8: 'ship', 9: 'truck'}
obb: None
orig_img: array([[[ 59,  78,  56],
                  [ 53,  71,  50],
                  [ 54,  72,  54],
                  ...,
                  [ 70,  87,  94],
                  [ 73,  87,  94],
                  [ 60,  73,  78]],

                 [[ 55,  72,  51],
                  [ 49,  65,  48],
                  [ 44,  57,  47],
                  ...,
                  [ 61,  81,  84],
                  [ 59,  79,  84]]])

```

```

[ 60,  73,  78]],

[[ 55,  72,  51],
 [ 49,  65,  48],
 [ 44,  57,  47],
 ...,
 [ 61,  81,  84],
 [ 62,  79,  84],
 [ 74,  89,  96]],

[[ 59,  72,  58],
 [ 54,  67,  55],
 [ 49,  60,  51],
 ...,
 [ 59,  81,  82],
 [ 66,  85,  89],
 [ 88, 103, 111]],

...,

[[ 73,  91,  94],
 [ 81,  97, 103],
 [ 75,  89,  97],
 ...,
 [ 75, 100,  90],
 [ 81, 111,  90],
 [ 79, 112,  91]],

[[ 82, 105,  97],
 [ 77,  97,  97],
 [ 80,  97, 104],
 ...,
 [ 76, 107,  99],
 [ 74, 107,  87],
 [ 68, 101,  78]],

[[ 60,  78,  72],
 [ 44,  59,  62],
 [ 47,  59,  67],
 ...,
 [ 82, 109, 115],
 [ 58,  83,  80],
 [ 40,  61,  57]]], dtype=uint8)
orig_shape: (32, 32)
path: '/content/datasets/cifar10/test/deer/0042.png'
probs: ultralytics.engine.results.Probs object
save_dir: '/content/runs/classify/predict'
speed: {'preprocess': 4.062497000177245, 'inference': 4.027097000289359, 'postprocess': 0.07067600017762743}]

```

The trained YOLOv8 model successfully classified sample images, demonstrating its effectiveness for image classification tasks.

Observations and Learnings

This assignment provided practical exposure to modern deep learning workflows and highlighted the importance of dataset structure and compatibility in machine learning pipelines. YOLOv8 offers a simple and efficient framework for image classification, making it suitable for rapid experimentation and deployment.

